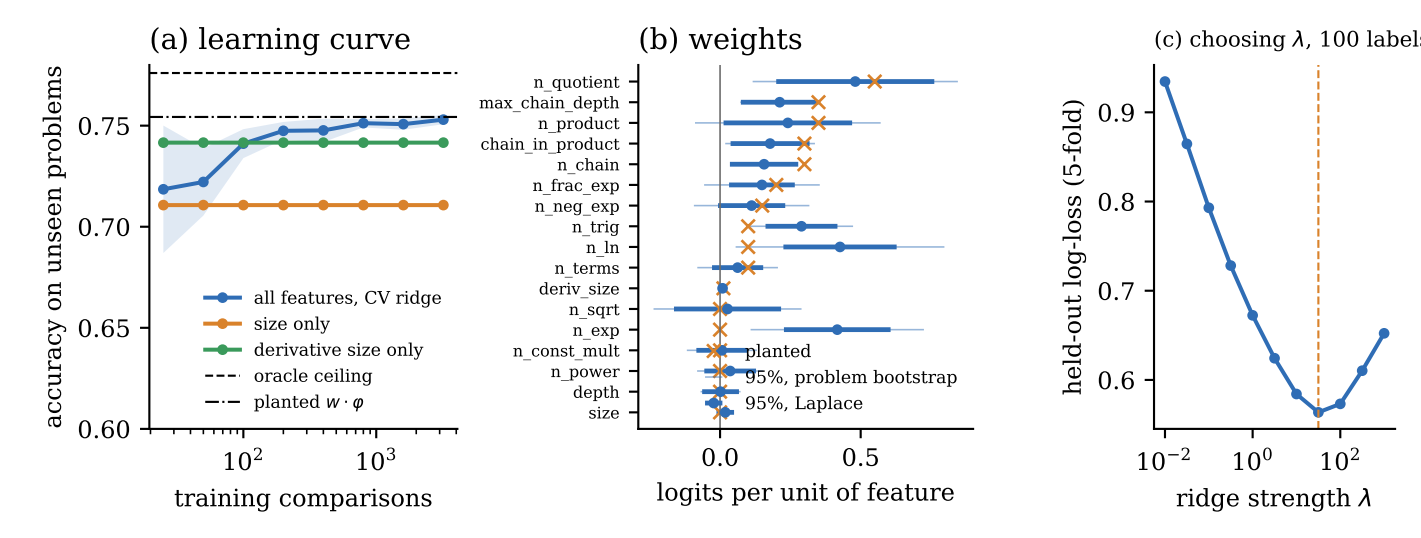
(a) learning curve
accuracy on unseen problems
0.75
0.70
0.65
0.60
10^2
10^3
training comparisons
all features, CV ridge
size only
derivative size only
oracle ceiling
planted $w \cdot \varphi$
(b) weights
n_quotient
max_chain_depth
n_product
chain_in_product
n_chain
n_frac_exp
n_neg_exp
n_trig
n_ln
n_terms
deriv_size
n_sqrt
n_exp
n_const_mult
n_power
depth
size
planted
95%, problem bootstrap
95%, Laplace
0.0
0.5
logits per unit of feature
(c) choosing λ, 100 labels
held-out log-loss (5-fold)
0.9
0.8
0.7
0.6
10^{-2}
10^0
10^2
ridge strength λ